

Mutexes and Monitors

Daniel Zappala

CS 360 Internet Programming
Brigham Young University

Mutexes

Mutex

- lock that allows only one thread into a critical section

```
1 #include <pthread.h>
2
3 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
4
5 int pthread_mutex_lock(pthread_mutex_t *mutex);
6 int pthread_mutex_trylock(pthread_mutex_t *mutex);
7 int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- must initialize the mutex first
- `pthread_mutex_lock()` will block if mutex is already locked
- `pthread_mutex_trylock()` will return EBUSY if mutex is locked

Don't Use Busy Waiting!

Busy Waiting

```
1 while running {
2     c = NULL;
3     pthread_mutex_lock(&mutex);
4     if queue.not_empty() {
5         c = queue.dequeue();
6     }
7     pthread_mutex_unlock(&mutex);
8     if c {
9         /* handle connection */
10    }
11 }
```

- must busy wait until a connection is available
- wastes CPU time on a server that does not handle many connections

Condition Variables

Condition Variables

```
1 #include <pthread.h>
2 pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
3
4 int pthread_cond_wait(pthread_cond_t *cond,
5                       pthread_mutex_t *mutex);
6
7 int pthread_cond_signal(pthread_cond_t);
```

- must initialize the condition variable first
- `pthread_cond_wait()` will block until the condition is signaled; the thread now owns the mutex as well
- need a corresponding `pthread_cond_signal()` to wake up

Using Condition Variables

```
1  while running {
2      c = NULL;
3      pthread_mutex_lock(&mutex);
4      while queue.empty() {
5          pthread_cond_wait(&cond,&mutex);
6      }
7      c = queue.dequeue();
8      pthread_mutex_unlock(&mutex);
9      /* handle connection */
10 }
```

- process inserting into queue should signal condition when queue goes from empty to having at least one item
- **must re-check queue status when conditional wait returns**
- no guarantee that queue will be empty when you return

Timed Wait and Broadcast Signals

```
1 #include <pthread.h>
2
3 int pthread_cond_timedwait(pthread_cond_t *cond,
4                             pthread_mutex_t *mutex,
5                             const struct timespec *abstime);
6
7 int pthread_cond_broadcast(pthread_cond_t *cond);
```


- `pthread_cond_timedwait()` needs an absolute time; use `clock_gettime()` and add the length of time you want to wait
- `pthread_cond_broadcast()` wakes up all threads waiting for a signal

Monitors

Monitor

- difficult to get semaphores, mutexes, condition variables right
 - match wait and signal
 - put in right order
 - scattered throughout code
- **monitor**: programming language construct
 - equivalent functionality
 - easier to control
 - mutual exclusion constraints can be checked by the compiler
 - used in versions of Pascal, Modula, Mesa
 - Java also has a Monitor object but compliance cannot be checked at compile time

Hoare Monitor



Hoare Monitor

- monitor can only be entered through methods
- shared memory can only be accessed by methods
- only one process or thread in monitor at any time
- may suspend and wait on a condition variable
- like object-oriented programming with mutual exclusion added in

Hoare Synchronization

- `cwait(c)`: suspend on condition `c`
- `csignal(c)`: wake up one thread waiting for condition `c`
 - do nothing if no threads waiting (signal is lost)
 - different from semaphore (number of signals represented in semaphore value)

Producer Consumer with a Hoare Monitor

```
1 vector buffer;  
2 condition notfull, notempty;
```

```
1 append(item) {  
2     if buffer.full()  
3         cwait(notfull);  
4     buffer.append(item);  
5     csignal(notempty);  
6 }
```

```
1 take() {  
2     if buffer.empty();  
3         cwait(notempty);  
4     item = buffer.remove();  
5     csignal(notfull);  
6     return item;
```

Producer Consumer with a Hoare Monitor

producer:

```
1 while (True) {  
2   item = produce();  
3   append(item);  
4 }
```

consume:

```
1 while (True) {  
2   item = take();  
3   consume(item);  
4 }
```

- advantages
 - moves all synchronization code into the monitor
 - monitor handles mutual exclusion
 - programmer handles synchronization (buffer full or empty)
 - synchronization is confined to monitor, so it is easier to check for correctness
 - write a correct monitor, any thread can use it

Lampson and Redell Monitor

- Hoare monitor requires that signaled thread must run immediately
 - thread that calls `csignal()` must exit the monitor or be suspended
 - for example, when `notempty` condition signaled, thread waiting must be activated immediately or else the condition may no longer be true when it is activated
 - usually restrict `csignal()` to be the last instruction in a method (Concurrent Pascal)
- Lampson and Redell
 - replace `csignal()` with `cnotify()`
 - `cnotify(x)` signals the condition variable, but thread may continue
 - thread at head of condition queue will run at some future time
 - must recheck the condition!
 - used in Mesa, Modula-3

Producer Consumer with a Lampson Redell Monitor

```
1 vector buffer;  
2 condition notfull, notempty;
```

```
1 append() {  
2     while buffer.full()  
3         cwait(notfull);  
4     buffer.append(item);  
5     cnotify(notempty);  
6 }
```

```
1 take() {  
2     while buffer.empty()  
3         cwait(notempty);  
4     item = buffer.remove();  
5     cnotify(notfull);  
6     return item;  
7 }
```

Lampson Redell Advantages

- allows processes in waiting queue to awaken periodically and reenter monitor, recheck condition
 - prevents starvation
- can also add `cbroadcast(x)`: wake up all processes waiting for condition
 - for example, append variable block of data, consumer consumes variable amount
 - for example, memory manager that frees k bytes, wake all to see who can go with k more bytes
- less prone to error
 - process always checks condition before doing work

What Can You Do?

- emulate a Lampson Redell Monitor with semaphores
 - create a class with private data only
 - use the same semaphore to protect all class methods
 - use semaphores to replace `cwait()` and `cnotify()`
- this creates a thread-safe class