

# Semaphores


Daniel Zappala

CS 360 Internet Programming  
Brigham Young University


# Semaphores

- **semaphore** is a shared variable maintained by OS
  - contains an integer and a queue
  - value initialized  $\geq 0$
- **wait(s)**: wait for a signal on semaphore *s*
  - decrements semaphore, blocks if value  $< 0$
  - if blocked, process put on the queue, suspends until signal is sent
- **signal(s)**: transmit a signal to semaphore *s*
  - increments semaphore
  - if value  $\leq 0$  then unblock someone
- **wait()** and **signal()** are atomic operations and cannot be interrupted

# wait()



# signal()



# Types of Sempahores

- **binary semaphore**
  - only one process at a time may be in the critical section
- **counting semaphore**
  - a fixed number of processes  $> 0$  may be in the critical section
- OS determines whether processes are released from queue in FIFO order or otherwise; usually FIFO in order to prevent starvation

# Using Semaphores

---

```
1 semaphore s = 1;
2
3 void thread(int i) {
4     while (true) {
5         wait(s);
6         /* critical section */
7         signal(s);
8         /* remainder */
9     }
10 }
```

---

- semaphore protects critical section
- can set  $s$  to  $> 1$  to let more than one process in the critical section
  - $s \geq 0$  : number that can enter
  - $s < 0$  : number that are waiting

# POSIX Semaphores

# POSIX Semaphores

---

```
1 #include <semaphore.h>
2
3 int sem_init(sem_t *sem, int pshared, unsigned int value);
4 int sem_wait(sem_t * sem);
5 int sem_trywait(sem_t * sem);
6 int sem_post(sem_t * sem);
```

---

- `sem_init()`: sets initial value of semaphore; `pshared = 0` indicates semaphore is local to the process
- `sem_wait()`: suspends process until semaphore is  $> 0$ , then decrements semaphore
- `sem_trywait()`: returns `EAGAIN` if semaphore count is  $= 0$
- `sem_post()`: increments semaphore, may cause another thread to wake from `sem_wait()`



# Example Code

- see example code `semaphore.cc`

► [GitHub](#)

**Producer Consumer**

# Producer Consumer Problem

- one or more producers are generating data and placing them in a buffer
- one or more consumers are taking items out of the buffer
- only one producer or consumer may access the buffer at any time

# Producer Consumer

---

```
1 vector buffer;  
2 append(item) {  
3     buffer.append(item);  
4 }  
5 take() {  
6     return buffer.remove();  
7 }
```

---

*producer:*

---

```
1 while (true) {  
2     item = produce();  
3     append(item);  
4 }
```

---

*consumer:*

---

```
1 while (True) {  
2     item = take();  
3     consume(item);  
4 }
```

---

# Producer Consumer with Infinite Buffer

---

```
1 sem_t s, n;  
2 sem_init(&s, 0, 1);  
3 sem_init(&n, 0, 0);
```

---

*producer:*

---

```
1 while (True) {  
2     item = produce();  
3     sem_wait(&s);  
4     append(item);  
5     sem_post(&s);  
6     sem_post(&n);  
7 }
```

---

*consumer:*

---

```
1 while (True) {  
2     sem_wait(&n);  
3     sem_wait(&s);  
4     item = take();  
5     sem_post(&s);  
6     consume(item);  
7 }
```

---

# Looking at the Code ...

- ① *What is the purpose of semaphore  $s$ ?*
- ② *What is the purpose of semaphore  $n$ ?*
- ③ *Why is semaphore  $s$  initialized to 1 but semaphore  $n$  is initialized to 0?*
- ④ *Why can the producer signal  $n$  every time an item is added to the buffer ?*
- ⑤ *Can the producer swap the signals for  $n$  and  $s$ ?*
- ⑥ *Can the consumer swap the waits for  $n$  and  $s$ ?*

# Important Insights

- two purposes for semaphores
  - **mutual exclusion**: semaphore  $s$  controls access to critical section
  - **signalling**: semaphore  $n$  coordinates when the buffer is empty: consumer waits if buffer is empty, producer signals when buffer becomes non-empty
- avoid race conditions
  - *item* keeps a local copy of the data protected by the semaphore so that it can be accessed later
  - reduces amount of processing inside the critical section

# Important Insights

- $n$ : semaphore value is number of items in buffer
  - if  $n == 0$ , consumer must wait
  - can swap `sem_post(&n);` and `sem_post(&s);` in producer and be OK
  - can't swap `sem_wait(&n);` and `sem_wait(&s);` in consumer: otherwise consumer enters and then waits and deadlocks the producer!
- ordering of semaphore operations is important



# Producer Consumer with Finite Buffer

---

```
1 sem_t s, n, e;  
2 sem_init(&s, 0, 1);  
3 sem_init(&n, 0, 0);  
4 sem_init(&e, 0, BUFFER_SIZE);
```

---

*producer:*

---

```
1 while (True) {  
2     produce();  
3     sem_wait(&e);  
4     sem_wait(&s);  
5     append();  
6     sem_post(&s);  
7     sem_post(&n);  
8 }
```

---

*consumer:*

---

```
1 while (True) {  
2     sem_wait(&n);  
3     sem_wait(&s);  
4     take();  
5     sem_post(&s);  
6     sem_post(&e);  
7     consume();  
8 }
```

---

# Looking at the Code ...

- 1 *What is the difference between semaphore e and semaphore n?*

# Thread-Safe Classes

# Organizing Semaphores

- difficult to get semaphores right
  - match wait and signal
  - put in right order
  - scattered throughout code
- put them in a class, with the data structures they use
  - private data structures, public methods
  - any object calling this class is thread-safe

# Thread-Safe Classes

---

```
1  class Buffer {
2      public:
3          append(item) {
4              sem_wait(&e);
5              sem_wait(&s);
6              buffer.append(item);
7              sem_post(&s);
8              sem_post(&n);
9          };
10         take() {
11             sem_wait(&n);
12             sem_wait(&s);
13             item = buffer.remove();
14             sem_post(&s);
15             sem_post(&e);
16             return item;
17         };
18
19     private:
20         vector buffer;
21 };
```